Strace in XV6

Xiang Mei
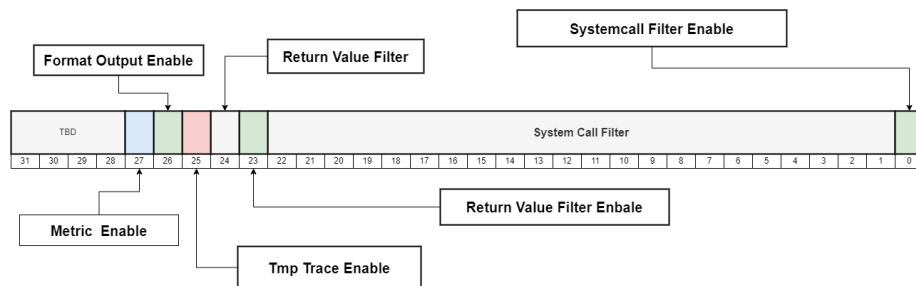xm2146@nyu.edu

April 28, 2022

1 Introduction

Figure 1: My Strace

Strace is a very useful tool on Linux. It's widely used to do the troubleshooting. But we don't have pre-installed strace on XV6. I'll implement a simple strace on XV6.

It sounds like reinventing a wheel. But for my expeirence in this course, Intro to OS, I think write a simple version of "wheel" could help me to understand the complesity of the "wheel" and help me think as an engineer. We need to consider about lots of questions during the implementation, such as "why should the wheel be round?". Anyway, I learned a lot and used some skills I learn from previous assignment.

The first section is the introduction of the report and I'll document my design in the second section. The real task-related parts start from the section 3.

2 Get familiar with Linux strace

In order to get familiar with the real strace, I use it to trace the sleep command. After reading the help page of strace, I use the "-C" flag to show the list of called syscalls, total number of calls, time of running strace on command.

```
[0] % strace -C sleep 1
execve("/bin/sleep", ["sleep", "1"], 0x7fff970a08f0 /* 46 vars */) = 0
brk(NULL)                               = 0x55932367a000
access("/etc/ld.so.preload", R_OK)      = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=34804, ...}) = 0
mmap(NULL, 34804, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f306c004000
close(3)                                = 0
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0@n\2\0\0\0\0\0"..., 832) = 832
fstat(3, {st_mode=S_IFREG|0755, st_size=1839792, ...}) = 0
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f306c002000
mmap(NULL, 1852680, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f306be3d000
mprotect(0x7f306be62000, 1662976, PROT_NONE) = 0
mmap(0x7f306be62000, 1355776, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x25000) = 0x7f306be62000
mmap(0x7f306bfad000, 303104, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x170000) = 0x7f306bfad000
mmap(0x7f306bff8000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1ba000) = 0x7f306bff8000
mmap(0x7f306bffe000, 13576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7f306bffe000
close(3)                                = 0
arch_prctl(ARCH_SET_FS, 0x7f306c003580) = 0
mprotect(0x7f306bff8000, 12288, PROT_READ) = 0
mprotect(0x5593219d8000, 4096, PROT_READ) = 0
mprotect(0x7f306c037000, 4096, PROT_READ) = 0
munmap(0x7f306c004000, 34804)           = 0
brk(NULL)                               = 0x55932367a000
brk(0x55932369b000)                     = 0x55932369b000
openat(AT_FDCWD, "/usr/lib/locale/locale-archive", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=3041456, ...}) = 0
mmap(NULL, 3041456, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f306bb56000
close(3)                                = 0
clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=1, tv_nsec=0}, 0x7ffe9d19ebe0) = 0
close(1)                                = 0
close(2)                                = 0
exit_group(0)                           = ?
+++ exited with 0 +++
% time     seconds  usecs/call     calls    errors syscall
------ ----------- ----------- --------- --------- ----------------
 53.73    0.000036          36         1           clock_nanosleep
 46.27    0.000031           6         5           close
  0.00    0.000000           0         1           read
  0.00    0.000000           0         3           fstat
  0.00    0.000000           0         8           mmap
  0.00    0.000000           0         4           mprotect
  0.00    0.000000           0         1           munmap
  0.00    0.000000           0         3           brk
  0.00    0.000000           0         1         1 access
  0.00    0.000000           0         1           execve
  0.00    0.000000           0         1           arch_prctl
  0.00    0.000000           0         3           openat
------ ----------- ----------- --------- --------- ----------------
100.00    0.000067           2        32         1 total
anubis@anubis-ide < master@1cb4ce3 > : ~/final-project-371525eb-xm2146 $
[0] %
```

Figure 2: Task 1-1

For task 1-2, I choose "mmap, read, execve, mprotect" as the targets to explain. We can see the procedure clearly on above figure. The "execve" call is called once and that's the first syscall we called. The "sleep" is an executable file in our system and the execve syscall could run it. It's worthy to mention the executed binary would use the caller's memory space and proc struct. By the way, the execve syscall has three 3 parameters, the first one is the path of the executable binary. And the second one would store the arguments while the third one would store the enviroment parameters.

The mmap syscall is used to allocate a large chunk of memory. In our command, it's used to allocate memory to store the shared libraries, the linker and other needed files. As we can see in above figure, mmap is usually called after openat. The first parameter of mmap syscall is the address of the allocated chunk. You can set it NULL to represent arbitrary address and we can also set it a non-NULL value to get a chunk strat from that adddress. This is used to allocate more space based on known chunk. The second parameter is the length we want to allocate while the thrid arguments is the

2

permission of the chunk, such as readable, writeable, and executeable.

And the mprotect is used to change the permission of memory. The mprotect can't allocate new memory. It could only change the permission of the memory chunks. In our command, it's used to make mmaped memory not writeable. For example, we allocate a chunk for the shared libraries and the memory must be writeable because we need copy the bytecodes to the memory. But you know it's dangerous to make writeable memory executable. So we need mprotect to make it un-writeable after copying.
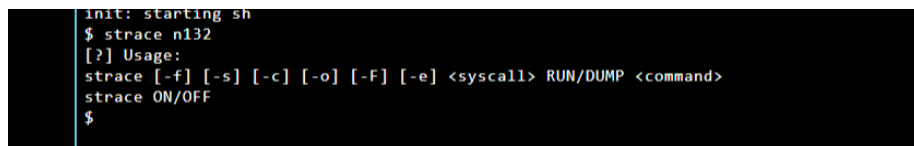
The syscall read is kind of straitforward. It reads the content from the first parameter's corresponding file and store to the second parameter's correctly memory while the third parameter is the max length of content the read syscall could read. It's used once to read the header of our glibc.

So far, we go through the usage and the shown information of strace on linux. Strace is a useful tools and I have been using it for a long time but I still find something new by reading it's help page. And we are going to implement out strace.

3 What features do we need

In this assignment, we gonna implement a simple strace. I'll go through and features needed.

For my implementation, the command should be like:


```
init: starting sh
$ strace n132
[?] Usage:
strace [-f] [-s] [-c] [-o] [-F] [-e] <syscall> RUN/DUMP <command>
strace ON/OFF
$
```

Figure 3: Usage

We have 4 sub-commands("RUN", "DUMP", "ON", and "OFF"), 3 filter options("-e","-s",and "-f"), and 3 output control options("-F","-c", and "-o").

I'll introduce these options and sub-commands in later section. And there are a short version introduction: The sub-command "RUN" could strace the following command and trace the syscall until the following command exits and the "DUMP" would dump the syscall records in the kernel. Also we can use "strace on" to ask the kernel keeps recording syscalls while the "strace off" could get the kernel back to the normal mode.

If we add "-e" options, the strace would only record and print specific syscalls. Besides, "-s/f" flag would force the strace to only record and print successful/failed syscalls.

Moreover, we have output control options. The "-F" flag would print more readable output and the "-c" would print a table of syscall metrics.

3

As for "-o", it can set an output file and the strace's output would be stored in that file.
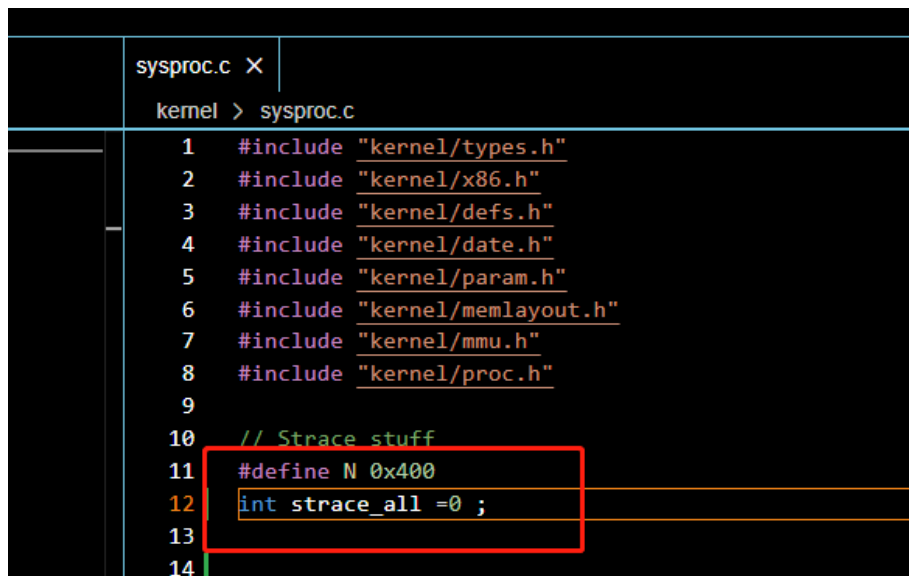
4   Design

This section will tell you the reasons of my design. This time, not like the "uniq" assignment, our implementation is different from the real sample because we don't have pthread syscall in xv6. Also, the strace is a big project I don't have much time to read its code. So during the implementation, I try to solve the problem myself and look up the matrials when I don't have an elegant solution.

ON/OFF

"When typing 'strace on' in the terminal, the mode of strace is on and therefore the next type in command will be traced. The system call list will be printed on screen in format pid (process id), command name, system call name, return value."

The "ON/OFF" implementation is easy and straitforward. In the requirements, we need to print the pid, name, and syscall of processes. Obviously, we can't do this task in userspace. I create an global variable in kernel space to present current strace_mode.
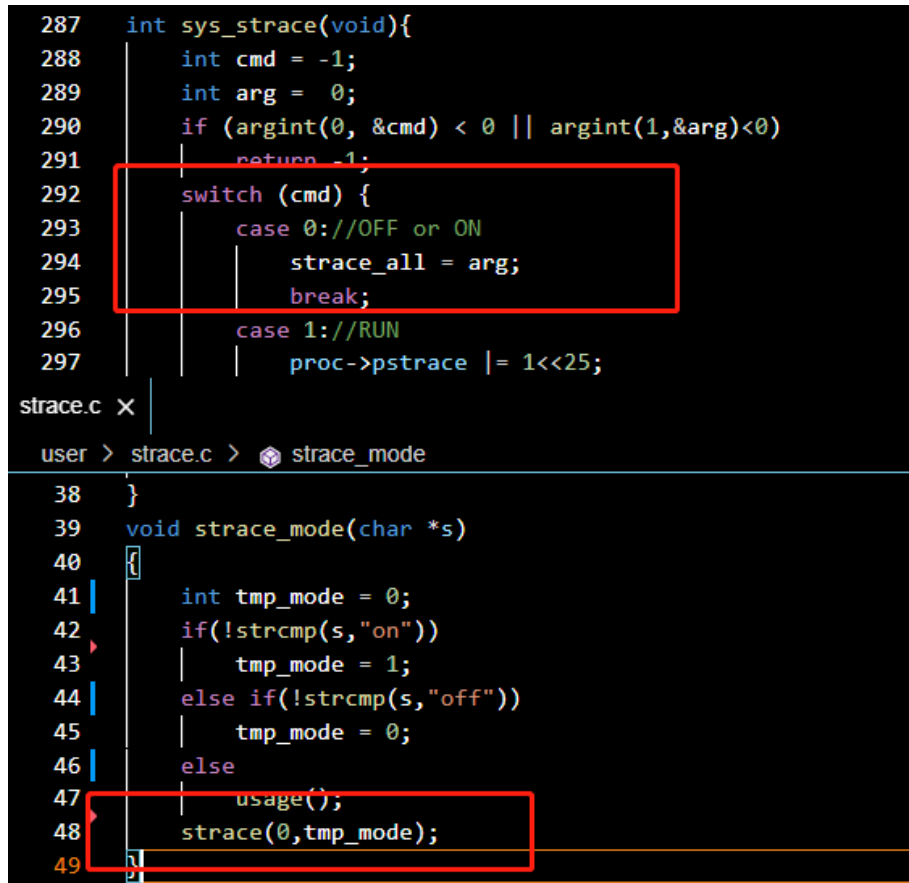
Figure 4: Strace All

Like what shows in the above figure, I set a global variable in "sysproc.c"

because I'll later implement a syscall as a bridge to stransfer command from
user space to the kernel.

```c
287    int sys_strace(void){
288        int cmd = -1;
289        int arg =  0;
290        if (argint(0, &cmd) < 0 || argint(1,&arg)<0)
291            return -1;
292        switch (cmd) {
293            case 0://OFF or ON
294                strace_all = arg;
295                break;
296            case 1://RUN
297                proc->pstrace |= 1<<25;
```

strace.c ✕

user > strace.c > ⬡ strace_mode

```c
38    }
39    void strace_mode(char *s)
40    {
41        int tmp_mode = 0;
42        if(!strcmp(s,"on"))
43            tmp_mode = 1;
44        else if(!strcmp(s,"off"))
45            tmp_mode = 0;
46        else
47            usage();
48        strace(0,tmp_mode);
49    }
```

Figure 5: Strace Syscall

As you can see in the above figure, I parse the parameters from the
command line and use strace syscall to modify the kernel variable. Basically,
we can use the code above to control the strace mode from user space.
Nevertheless, how do we monitor the syscalls.

At very first, a simple plan came up in my mind: I can insert a piece
of code in every syscall so that I can print the result and parameters when
running the syscall. And I use a global variable "strace_all" to tell the
system if we should print the strace info while calling syscalls. But quickly
I found there are some serious issues with this solution. We have 21 syscalls
and if we need to write different strace_handle for every syscall and it's not
convinience because we may need to modify 21 places for every single change.
My expeirence told me it's horriable. We must implement something more
elegant.

The advantage of the previous plan is that we can print the arguments

of syscalls. If we need to print the content of the syscall we have to do that because the syscalls would parse the parameters in these syscall handlers. I asked in slack and found we don't need to print the details about the syscall so that we can move the strace code to the syscall_interupt handle or the wrapper function.

Figure 6: Syscall Handler in Trap Fcuntion

As we talked in this trap section, the syscall in user space would use create an interupt to inform the kernel. And the interup is handled by the function alltraps in "trapasm.S" it would store current context in the trap frame and call function trap which is shown on the above figure. And we can see the trap would check the process's state and call the function syscall. This function is in file syscall.c.

Figure 7: Syscall Wrapper

In this function, we would parse the EAX which represent the syscall index and store the return value in the EAX. So I think this function is a good candidate to insert our strace code.

6

Figure 8: Monitor

I insert the monitor after the syscall because we need the return value the syscall but we may loss "exit" syscall because the process would stop in "exit" syscall. So I add the same function before the process really exits which is shown in the figure 9.

Figure 9: Strace Handler in sys_exit

The usage of this sub-command is simple. You can just use "strace on" to turn on strace and "strace off" to turn off strace.

Figure 10: Strace On and OFF

There is no thing more about these two sub-commands but there is little problem about the global variable. In order for avoiding race condition, we need to implement lock mechanism for this variable. Howerver, this not the

7

requirement for this assignment and there is only one member in my team. I decide to do implement that only if I have time left.

DUMP

"Implement a kernel memory that will save N number of latest events. This N number can be configurable by using define in XV6. In order word, it can be hard code but the way to implement is to use define to declare a variable called N with certain value. When 'strace dump' command is called, print all events that saved in kernel memory."

In previosu figures, you may noticed that I implemented a function to log the syscall. so why do I implement a such conplex function "add_one_record".

For implementing the "DUMP" feature, we need to allocate a space in the kernel to store the latest N system calls. These data have to be stored in the kernel space as a global variable because xv6 is a multi-process system. And we need a special circle buffer to store latest N records.

```
12      while(1) // Record Inputs
13      {
14          read(0,&c,1);
15          if(c==0xa) break;
16          buf[cur] = c;
17          cur++;
18          if(cur>=N){
19          cur = 0; flag=1;
20          }
21      }
22      //Dump Inputs
23      if(flag==0)
24          for(i = 0 ;  i < cur ;i++)
25              putchar(buf[i]);
26      else{
27          i = cur;
28          do{
29              putchar(buf[i]);
30              i++;
31              if(i>=N) i = 0;
32          }while(i!=cur);
33      }
34      putchar(0xa);
```

Figure 11: Circle Buffer

Circle buffer for DUMP operation is easier than general circle buffer. First, I implemented a C code version for testing. As you can see in the above figure, We read the input to the circle buffer and dump the content when we get an "enter". The read-part is simple and the "flag" variable is important. It decides how many and where to dump. The trick in the code is simple and makes sure we would print the last N records which is veried by my fuzzer. That's another reason why I write it in C. Also, I attach my

fuzz code:

```python
6  def payload_gen(l):
7      res = ''
8      for x in range(l):
9          res+=table[random.randint(0,len(table)-1)]
10     return res
11 def fuzzer():
12     p = process("./main")
13     pay = payload_gen(random.randint(1,0x100))
14     p.sendline(pay)
15     output = p.readline()[:-1].decode()
16     if output != pay[-N:]:
17         print("Bug Found")
18         exit(1)
19     p.close()
20 if __name__ == "__main__":
21     while(1):
22         fuzzer()
```

Figure 12: Circle Buffer Fuzzer

After other testing, I implement a similar circle buffer in the kernel to store and dump the syscalls. In function "add_one_record", I record the syscall's inform in the correct node and move the pointer like what I did in my previous demo.

10

```
188    void add_one_record(unsigned int pstrace,int pid,char *name,uint
189    {
190 >      if( !(pstrace&(1<<25)) && !strace_all)//Pruned Mode···
194 >      if(!strncmp(name, "strace", 7) || !strncmp(name, "sh", 3))//
196 >      if(pstrace&(1<<23)){···
201        }// -e or -c
202 >      if( pstrace & 1 ){···
205        }
206 >      else if ( pstrace & (1<<27)){···
211        }
212        strace_record[strace_cur].pid = pid;
213        strncpy(strace_record[strace_cur].name,name,0x10-1);
214        strace_record[strace_cur].sys_id = sys_id;
215        strace_record[strace_cur].ret_val = ret_val;
216        struct strace_node * p = &strace_record[strace_cur];
217 >      if(!(pstrace &(1<<26)))···
219        strace_cur++;
220        if(strace_cur >= N)
221        {
222            strace_cur = 0;
223            strace_flag = 1;
224        }
225        return;
226    }
```

Figure 13: Circle Buffer Fuzzer

That's the key function of all my design I'll metion this function later to introduce otehr features. This function is only called in function syscall and sys_exit so that if we want to modify, delete, or add a feature to strace, we can just modify the code in this function. Another advantage is that we can naturally combine kinds of options.

```
TRACE: pid = 3  |  Command_name = grep  |  syscall = exit  |
$ strace dump
TRACE: pid = 3  |  Command_name = grep  |  syscall = exec   |  Return value = 0
TRACE: pid = 3  |  Command_name = grep  |  syscall = open   |  Return value = 3
TRACE: pid = 3  |  Command_name = grep  |  syscall = read   |  Return value = 1023
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 371
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 34
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 114
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 120
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 121
TRACE: pid = 3  |  Command_name = grep  |  syscall = read   |  Return value = 853
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 43
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 193
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 68
TRACE: pid = 3  |  Command_name = grep  |  syscall = read   |  Return value = 991
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 45
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 71
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 47
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 76
TRACE: pid = 3  |  Command_name = grep  |  syscall = read   |  Return value = 811
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 45
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 78
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 67
TRACE: pid = 3  |  Command_name = grep  |  syscall = write  |  Return value = 66
TRACE: pid = 3  |  Command_name = grep  |  syscall = read   |  Return value = 0
TRACE: pid = 3  |  Command_name = grep  |  syscall = close  |  Return value = 0
TRACE: pid = 3  |  Command_name = grep  |  syscall = exit   |
$
```

Figure 14: DUMP after Traing "grep the readme.md"

I did several tests on DUMP and it works well. I attach a simple sample above because the output of complex testcase would be big and hard to recognize. You can also test it with any commands you like and please check the README.md file attached to the assignment submission.

RUN

> "Instead of turning on and off strace, we create 'strace run' to directly point tracing to the current process that execute the command. For example: when typing 'strace run echo hello' in the terminal, we get the output tracing of echo hello."

For sub-command "RUN", we gonna run a command and strace the syscall of the command. So we can parse the parameter and use "fork" or "exec" syscall to run the command.

I used the "exec" syscall for my strace and there is a disadvantage comparing with the fork version. I know the fork version from my professor: we can use "fork" to creat a child process and let the parent process wait for the metadata of outputs from the kernel. So all the output part would be handledin user space which is much more secure and beautiful. Howerver, I almost finish my implementation and there is no much time for this huge modification. Therefore, I'll keep my "exec" version and for my implementation I think they are similar on complesity.
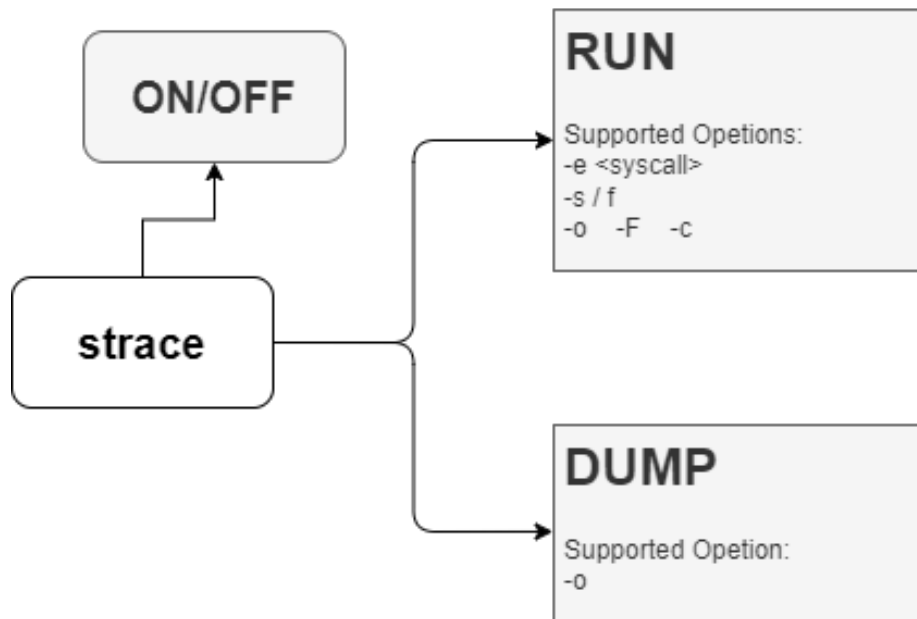
Figure 15: Supported Opetions

The "RUN" sub-command is the most complex part of my strace, I'll introduce some basic idea about this sub-command in this parameter and explain the rest parts with the supported options. So how can we strace one command once?

If we "RUN" a command, we are going to use "fork" to create a new process or use "exec" to run a executable file by using current memory space. We must have a way to pass the information that the new process shoud be traced. In order to implement this, I add a new element in "proc" struct in proc.

Figure 16: Strace a Process

As you can see in the above figure, I add a unsigned interger to the proc struct. I don't want to wast unnecessary space in the kernel space. so I would use this 4 bytes variable to store all the strace information such as output filter information and output format information. I'll explain the struct of this variable in option-related paragraph.

Figure 17: RUN Sub-command in User Space

The above figure is my userspace interface and corresponding system call implementation. As you can see I use the 26th bit of the "pstrace" to sign if the kernel should strace the process. Another advantage of having a variable in the "proc" struc is that we can easily follow the subprocess by modify the "fork" function in "proc.c".

Figure 18: Trace the Child Processes

Now we can use strace to run strace any process! There is a simple demo of "RUN" sub-command.

Figure 19: Strace RUN

Options

My strace supports kinds of output filter and format options which could help to eliminate the uninterested sycalls.

Figure 20: Supported Opetions

I would parse the parameters in the user space and use the strace syscall to pass the operations to the kernel mode. As we talked in previous paragraph, the atomic unit of strace is the process so we need a variable in proc stuct to tell strace-kernel what shoud it do. I'll introduce every bits of "pstrace" variable in following paragraphs. And you can check the struct of the variable in the following figure:
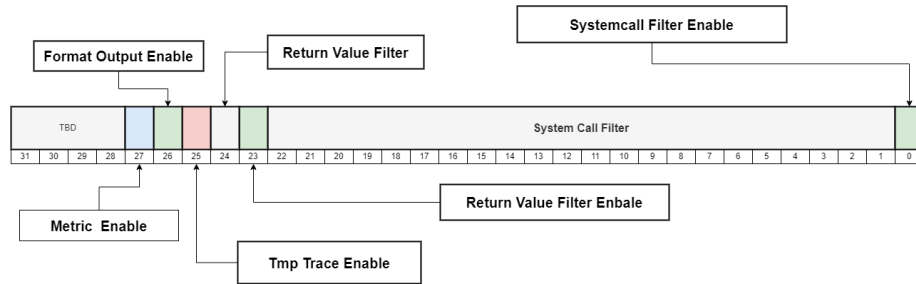

Figure 21: "pstrace" in "proc"

Option -e [system call name]: When option flag -e is provided follow by a system call name, we will print only that system call. If no such system call is made in the command, print nothing.

Option -s: When option flag -s is provided, print only successful system call.

Option -f: When option flag -f is provided, print only failed system call.

Option -s: Options -c in strace will generate a statistic report of system call regarding the input command such as duration, total call, failed call. Create a similar report table with using option -c.

Option -f: When option flag -f is provided, print more readable output.

We have 21 syscalls on xv6 and 5 different opetions, so we can store these informations in a 32-bit variable. As you can see in the figrue, I use 21+1 bit for "-e" option.

Option -e The 0th bit is the inuse-bit of 1-22 bits. For example, if we don't run strace with the "-e" flag the 0th bit is 0, the strace-kernel would not check the syscall filter.

Figure 22: -e Opetion Implementation

The above figure is the user space and the kernel space handler of "-e" operation. The user space handler would use strace syscall to pass the filter to the kernel space which is a whitelist of syscalls while the kernel space would enable the syscall filter and apply the filter.

Figure 23: -e Opetion Demo

The above figure is a screenshot of the usage of "-e" opetion and as you can see, my implementation could naturally handle combination of "-e" opetion.

Option -s/f The 23th and 24th bit of pstrace is design for "-s/f" flag. The 23th bit is the inuse-bit of 24th bit. And if 24th bit is 0, the kernel would only record successful syscalls while if 24th bit is 1, the kernel would only record failed syscalls.

17

```
71              cur++;
72              filter |=syscall2int(argv[cur]);
73              strace(4,filter);
74              break; //("e")
75          case 115:
76              strace(3,0);
77              break; //("s")
78          case 102:
79              strace(3,1);
80              break; //("f")
81          case 70:
82              strace(6,0);
```

sysproc.c ✕

kernel > sysproc.c > sys_strace

```
285              strace_dump(-1);
286              break;
287          case 3://Fail or success 23/24
288              proc->pstrace |= (1<<23);
289              // 0 means ret filter OFF(Default State)
290              // This step would trun it on(change the 23th bit to 1)
291              if(arg==1)
292                  proc->pstrace |= (1<<24);
293              else if(arg==0)
294                  proc->pstrace &= (~(1<<24));
295              else
296                  panic("SYS_strace");
297              break;
```
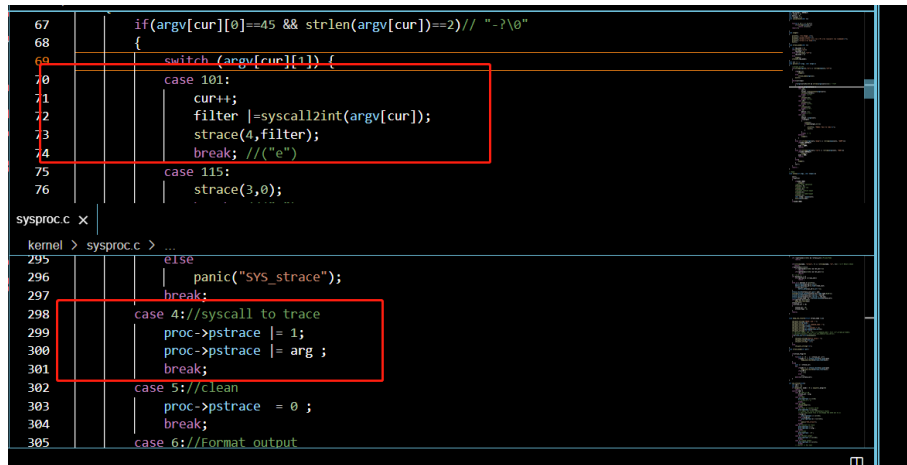
Figure 24: -s/f Opetion Implementation(Handler)

The above figure is the user space and the kernel space handler of "-s/f"
operation. The user part would simply parse the arguments and call the
kernel to apply the filter. Also, the kernel mode would apply the operation
to the pstrace of the proc. And as you can see in the beneath figure we will
check the filter before printing out the syscall in "add_one_record".

```
8    void add_one_record(unsigned int pstrace,int pid,char *name,
9  ∨ {
0  >      if( !(pstrace&(1<<25)) && !strace_all)//Pruned Mode···
2  >      if(!strncmp(name, "strace", 7) || !strncmp(name, "sh", 3
4  ∨      if(pstrace&(1<<23)){
5  ∨          if ((pstrace&(1<<24)) && ret_val!=-1)
6               return;
7  ∨          if(!(pstrace&(1<<24)) && ret_val==-1)
8               return;
9  |      }// -s / f
```

Figure 25: -s/f Opetion Implementation(Filter)

Figure 26: -s/f Opetion Demo

So far we finish the introduction of filter options (-s/f/e), and we'll move to format options in the following paragraphs.

Option -F

"Depends on your implementation, you might notice your command result is printed along with your strace result which cause tracing to be difficult to read when characters keep mix up between line (ex: 'ls' has a very long list result). Choose one of the follow approach to solve this issue:"

Figure 27: -F Opetion Implementation(Handler)

For this option, we need to "pause" the stace output until the process finished its output. So I would store the output in someplace and dump it after the process exits by reuse partial code of "DUMP" sub-command. This is quite simple and elegant but it has a little flaw: What if the process calss tons of syscall and we could only store latest N records.

```
2   int sys_exit(void) {
3   |
4   |    //if(!strncmp(proc->name, "strace", 7) || !strncmp(proc->name, "sh", 3))
5   |    //    exit();
6   |    add_one_record(proc->pstrace,proc->pid,proc->name,2,(int)0xdeadbeef,0);
7   |    if((proc->pstrace) & (1<<26))
8   |    {
9   |        if((proc->pstrace) & (1<<27))
0   |        {
1   |        |    metirc();
2   |        }
3   |        else
4   |        |    strace_dump(proc->pid);
5   |    }
6   |    exit();
7   |    return 0; // not reached
8   }
```

Figure 28: -F Opetion Implementation

For this issue, I can just set a variable to monitor the storage and dump the records if we are in "-F" mode. But I think this method would break XV6's simply so I would rather mention this issue in the README.md. And the follow figrue shows the output of "-F" option.

```
init: starting sh
$ strace -F run ls
.              1 1 512
..             1 1 512
README.md      2 2 3678
cat            2 3 15512
echo           2 4 14620
forktest       2 5 9232
grep           2 6 17620
init           2 7 15208
kill           2 8 14728
ln             2 9 14620
ls             2 10 17096
mkdir          2 11 14772
rm             2 12 14756
sh             2 13 27640
stressfs       2 14 15324
usertests      2 15 65184
wc             2 16 16240
zombie         2 17 14292
strace         2 18 19308
deb            2 19 14344
console        3 20 0
TRACE: pid = 3  |        Command_name = ls   |        syscall = ex
TRACE: pid = 3  |        Command_name = ls   |        syscall = op
TRACE: pid = 3  |        Command_name = ls   |        syscall = fs
```

Figure 29: -F Opetion Demo

Option -c

21

"Options -c in strace will generate a statistic report of system call regarding the input command such as duration, total call, failed call. Create a similar report table with using option -c."

The "-c" would show the metrics of the command. We need to store more informations abotu the syscalls so that allocating more space is necessary. For time recording, I would use uptime syscall to calculate the time used for every syscall. And it's easy to count the error number and usage number in function "add_one_record".

```
168    void add_one_record(unsigned int pstrace,int pid,char *name,ui
169    {
170  >     if( !(pstrace&(1<<25)) && !strace_all)//Pruned Mode ⋯
172  >     if(!strncmp(name, "strace", 7) || !strncmp(name, "sh", 3))
174  >     if(pstrace&(1<<23)){ ⋯
179        }// -s / f
180        if( pstrace & 1 ){
181            if(!(pstrace & (1<<sys_id)))
182                return ;
183        }
184        if ( pstrace & (1<<27)){
185            metric_meta[sys_id-1].ct+=1;
186            metric_meta[sys_id-1].time+=time_cost;
187            if(ret_val==-1)
188                metric_meta[sys_id-1].err +=1;
189        }
190        strace_record[strace_cur].pid = pid;
191        strncpy(strace_record[strace_cur].name,name,0x10-1);
```

Figure 30: -C Option Implementation

And just before the exit, the "metric" function would be triggered and dump the "-c" related informations as the following figure.

```
71
72    void metirc()
73    {
74        std_puts_string("time\t\tcalls\t\terrors\t\tsyscall\n");
75        std_puts_string("-----------------------------------------------\n");
76        for(int i=0;i<21;i++)
77        {
78            if(metric_meta[i].ct!=0)
79            {
80                std_puts_int(metric_meta[i].time);
81                std_puts_string("\t\t");
82                std_puts_int(metric_meta[i].ct);
83                std_puts_string("\t\t");
84                std_puts_int(metric_meta[i].err);
85                std_puts_string("\t\t");
86                std_puts_string(sys_name[i]);
87                std_puts_string("\n");
88            }
89        }
90        metric_clean();
91    }
```

Figure 31: -C Option Demo

```
$ strace -c run ls
.               1 1 512
..              1 1 512
README.md       2 2 3678
cat             2 3 15512
echo            2 4 14620
forktest        2 5 9232
grep            2 6 17620
init            2 7 15208
kill            2 8 14728
ln              2 9 14620
ls              2 10 17096
mkdir           2 11 14772
rm              2 12 14756
sh              2 13 27640
stressfs        2 14 15324
usertests       2 15 65184
wc              2 16 16240
zombie          2 17 14292
strace          2 18 19308
deb             2 19 14344
console         3 20 0
time            calls           errors          syscall
------------------------------------------------------------
0               1               0               exit
0               33              0               read
1               1               0               exec
0               22              0               fstat
1               22              0               open
2               526             0               write
0               22              0               close
$
```
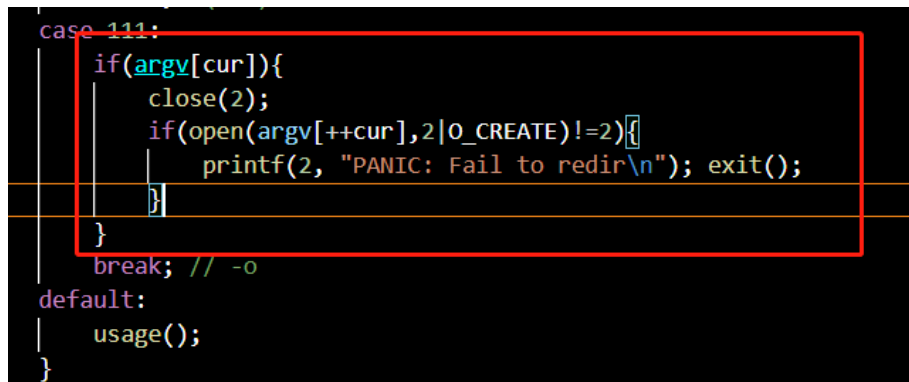
Figure 32: -C Option Demo

23

And it's great to know "evec" costs such much (33 times more than read)!

Option -o

> "Find an implementation of choice to write strace output to file. For example: by providing option: -o 'filename' or editing content of README."

Figure 33: -o Option Userspace Handler

This one is the hardests one for my implementation because I use "exec" to run the command rather than the fork. So I need to handle all the things in the kernel which is not secure. As you see in above figure, I can easily finish the userspace interface. Similar to what we learned in "sh.c", I just redirecte the stderr to the given file.

```
void std_puts_string(char *s)
{
    filewrite(proc->ofile[2], s, strlen(s));
}
void std_puts_int(int n)
{

    if(n<0){
        std_puts_string("-");
        n = -n;
    }
    uint idx = 0 ;
    char s[0x20]={0};
    //only 32 are used,
    //but it's better for alignment to use 32;
    do{
        s[idx++] = (char)((n%10)+0x30);
        n = n/10;
    }while(n!=0);
    //Out put
    while(idx!=0)
        filewrite(proc->ofile[2], &s[--idx], 1);
}
```

Figure 34: Format Print in Kernel

But unluckly, we don't have "printf" in the kernel to store the data in to stderr. So I need to implement a limited "printf" to store the output to the stdin out. I use readfile to achieve that as above code shows. I replace all output functions with my "format print" functions so that "-o" could naturally work with other commands and options, which is shown in the beneath figure.

```
init: starting sh
$ strace -o n132 run echo n132
n132
$ cat ./n132
TRACE: pid = 3    |        Command_name = echo    |        syscall = exec   |        Return value = 0
TRACE: pid = 3    |        Command_name = echo    |        syscall = write  |        Return value = 1
TRACE: pid = 3    |        Command_name = echo    |        syscall = write  |        Return value = 1
TRACE: pid = 3    |        Command_name = echo    |        syscall = write  |        Return value = 1
TRACE: pid = 3    |        Command_name = echo    |        syscall = write  |        Return value = 1
TRACE: pid = 3    |        Command_name = echo    |        syscall = write  |        Return value = 1
TRACE: pid = 3    |        Command_name = echo    |        syscall = exit   |
$ []
```

Figure 35: -o Opetion Demo

# 5 Usage Details

6 Testing

7 Todo

8 Summary

Figure 36: Structure

There are two main parts of this implementation: Kernel Space Part and User Space Part.

In the user space I use the binary strace as the interface to transfer data to the kernel. And the main handler is the file "kernel/proc.c". I implement an strace syscall and deal with kinds of parameters.

In the kernel space, I use the metadata got from the strace syscall to set the output format/filter options. Also, I use the kernel variables to implemt "strace on/off".

And you can check all the supported features in above figrue.